

Problem statement: To build a CNN based model which can accurately detect melanoma. Melanoma is a type of cancer that can be deadly if not detected early. It accounts for 75% of skin cancer deaths. A solution which can evaluate images and alert the dermatologists about the presence of melanoma has the potential to reduce a lot of manual effort needed in diagnosis.

Importing Skin Cancer Data

To do: Take necessary actions to read the data

Importing all the important libraries

In [146]:

```
1 import pathlib
2 import tensorflow as tf
3 import matplotlib.pyplot as plt
4 import numpy as np
5 import pandas as pd
6 import os
7 import PIL
8 from tensorflow import keras
9 from tensorflow.keras import layers
10 from tensorflow.keras.models import Sequential
11 from glob import glob
12 import warnings
13 warnings.filterwarnings('ignore')
```

In []:

```
1 ## If you are using the data by mounting the google drive, use the following :
2 # from google.colab import drive
3 # drive.mount('/content/gdrive')
4 #!unzip gdrive/MyDrive/CNN_assignment
5 ##Ref:https://towardsdatascience.com/downloading-datasets-into-google-drive-via-google-
6
7
8 %rm -rf 'Skin cancer ISIC The International Skin Imaging Collaboration'
9 !gdown --id 1xLfSQUGD18ezNNbUkpuHOYvSpTyxVhCs
10 !unzip -o -q CNN_assignment.zip
```

This assignment uses a dataset of about 2357 images of skin cancer types. The dataset contains 9 sub-directories in each train and test subdirectories. The 9 sub-directories contains the images of 9 skin cancer types respectively.

In [106]:

```
1 # Defining the path for train and test images
2 ## Todo: Update the paths of the train and test dataset
3 data_dir_train = pathlib.Path("/content/Skin cancer ISIC The International Skin Imaging")
4 data_dir_test = pathlib.Path("/content/Skin cancer ISIC The International Skin Imaging")
```

In [107]:

```
1 !ls
```

```
CNN_assignment.zip
model_plot.png
sample_data
'Skin cancer ISIC The International Skin Imaging Collaboration'
```

In [108]:

```
1 image_count_train = len(list(data_dir_train.glob('*/*.jpg')))
2 print(image_count_train)
3 image_count_test = len(list(data_dir_test.glob('*/*.jpg')))
4 print(image_count_test)
```

```
2239
```

```
118
```

Load using keras.preprocessing

Let's load these images off disk using the helpful `image_dataset_from_directory` utility.

Create a dataset

Define some parameters for the loader:

In [109]:

```
1 batch_size = 32
2 img_height = 180
3 img_width = 180
```

Use 80% of the images for training, and 20% for validation.

In [110]:

```
1 ## Write your train dataset here
2 ## Note use seed=123 while creating your dataset using tf.keras.preprocessing.image_data
3 ## Note, make sure you resize your images to the size img_height*img_width, while writ
4 train_ds = tf.keras.preprocessing.image_dataset_from_directory(
5     data_dir_train,
6     seed=123,
7     validation_split= 0.2,
8     subset= 'training',
9     image_size=(img_height,img_width),
10    batch_size = batch_size
11 )
```

Found 2239 files belonging to 9 classes.
Using 1792 files for training.

In [111]:

```
1  ## Write your validation dataset here
2  ## Note use seed=123 while creating your dataset using tf.keras.preprocessing.image_data_utils
3  ## Note, make sure you resize your images to the size img_height*img_width, while writing
4  val_ds = tf.keras.preprocessing.image_dataset_from_directory(
5      data_dir_train,
6      seed=123,
7      validation_split= 0.2,
8      subset= 'validation',
9      image_size=(img_height,img_width),
10     batch_size = batch_size
11 )
```

Found 2239 files belonging to 9 classes.
Using 447 files for validation.

In [112]:

```
1  # List out all the classes of skin cancer and store them in a list.
2  # You can find the class names in the class_names attribute on these datasets.
3  # These correspond to the directory names in alphabetical order.
4  class_names = train_ds.class_names
5  class_names
```

Out[112]:

```
['actinic keratosis',
 'basal cell carcinoma',
 'dermatofibroma',
 'melanoma',
 'nevus',
 'pigmented benign keratosis',
 'seborrheic keratosis',
 'squamous cell carcinoma',
 'vascular lesion']
```

Visualize the data

Todo, create a code to visualize one instance of all the nine classes present in the dataset


In [113]:

```
1 ### your code goes here, you can use training or validation data to visualize
2 import matplotlib.pyplot as plt
3
4 plt.figure(figsize=(15, 15))
5 for images, labels in train_ds.take(1):
6     for i in range(9):
7         ax = plt.subplot(3, 3, i + 1)
8         plt.imshow(images[i].numpy().astype("uint8"))
9         plt.title(class_names[labels[i]])
10        plt.axis("off")
```


vascular lesion




nevus




pigmented benign keratosis




pigmented benign keratosis




melanoma




dermatofibroma




pigmented benign keratosis



actinic keratosis



basal cell carcinoma



The `image_batch` is a tensor of the shape `(32, 180, 180, 3)`. This is a batch of 32 images of shape `180x180x3` (the last dimension refers to color channels RGB). The `label_batch` is a tensor of the shape `(32,)`, these are corresponding labels to the 32 images.

`Dataset.cache()` keeps the images in memory after they're loaded off disk during the first epoch.

`Dataset.prefetch()` overlaps data preprocessing and model execution while training.

In [114]:

```
1 AUTOTUNE = tf.data.experimental.AUTOTUNE
2 train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
3 val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
```

Create the model

Todo: Create a CNN model, which can accurately detect 9 classes present in the dataset. Use `layers.experimental.preprocessing.Rescaling` to normalize pixel values between (0,1). The RGB channel values are in the `[0, 255]` range. This is not ideal for a neural network. Here, it is good to standardize values to be in the `[0, 1]`

In [115]:

```
1 ### Your code goes here
2 from keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPool2D
3 num_classes = 9
4 model = Sequential([
5     layers.experimental.preprocessing.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
6     layers.Conv2D(16, 3, padding='same', activation='relu'),
7     layers.MaxPooling2D(),
8     layers.Conv2D(32, 3, padding='same', activation='relu'),
9     layers.MaxPooling2D(),
10    layers.Conv2D(64, 3, padding='same', activation='relu'),
11    layers.MaxPooling2D(),
12    layers.Flatten(),
13    layers.Dense(128, activation='relu'),
14    layers.Dense(num_classes)
15 ])
```

Compile the model

Choose an appropriate optimiser and loss function for model training

In [116]:

```
1 ### Todo, choose an appropriate optimiser and loss function
2 model.compile(optimizer='adam',
3               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
4               metrics=['accuracy'])
```

In [117]:

```
1 # View the summary of all layers
2 model.summary()
```

Model: "sequential_21"

Layer (type)	Output Shape	Param #
=====		
rescaling_19 (Rescaling)	(None, 180, 180, 3)	0
conv2d_57 (Conv2D)	(None, 180, 180, 16)	448
max_pooling2d_57 (MaxPooling2D)	(None, 90, 90, 16)	0
conv2d_58 (Conv2D)	(None, 90, 90, 32)	4640
max_pooling2d_58 (MaxPooling2D)	(None, 45, 45, 32)	0
conv2d_59 (Conv2D)	(None, 45, 45, 64)	18496
max_pooling2d_59 (MaxPooling2D)	(None, 22, 22, 64)	0
flatten_19 (Flatten)	(None, 30976)	0
dense_44 (Dense)	(None, 128)	3965056
dense_45 (Dense)	(None, 9)	1161
=====		
Total params: 3,989,801		
Trainable params: 3,989,801		
Non-trainable params: 0		

Train the model

In [118]:

```
1 epochs = 20
2 history = model.fit(
3     train_ds,
4     validation_data=val_ds,
5     epochs=epochs
6 )
```

Epoch 1/20

56/56 [=====] - 14s 70ms/step - loss: 1.8175 - accuracy: 0.3158 - val_loss: 1.6135 - val_accuracy: 0.4430

Epoch 2/20

56/56 [=====] - 1s 24ms/step - loss: 1.5158 - accuracy: 0.4648 - val_loss: 1.5130 - val_accuracy: 0.4922

Epoch 3/20

56/56 [=====] - 1s 23ms/step - loss: 1.3949 - accuracy: 0.5089 - val_loss: 1.3946 - val_accuracy: 0.4944

Epoch 4/20

56/56 [=====] - 1s 23ms/step - loss: 1.2901 - accuracy: 0.5474 - val_loss: 1.3601 - val_accuracy: 0.5190

Epoch 5/20

56/56 [=====] - 1s 24ms/step - loss: 1.1835 - accuracy: 0.5820 - val_loss: 1.4396 - val_accuracy: 0.5034

Epoch 6/20

56/56 [=====] - 1s 24ms/step - loss: 1.0999 - accuracy: 0.6088 - val_loss: 1.3980 - val_accuracy: 0.5369

Epoch 7/20

56/56 [=====] - 1s 24ms/step - loss: 1.0016 - accuracy: 0.6323 - val_loss: 1.4212 - val_accuracy: 0.5324

Epoch 8/20

56/56 [=====] - 1s 24ms/step - loss: 0.9411 - accuracy: 0.6551 - val_loss: 1.3582 - val_accuracy: 0.5593

Epoch 9/20

56/56 [=====] - 1s 24ms/step - loss: 0.8231 - accuracy: 0.7048 - val_loss: 1.4639 - val_accuracy: 0.5414

Epoch 10/20

56/56 [=====] - 1s 24ms/step - loss: 0.7954 - accuracy: 0.7204 - val_loss: 1.5935 - val_accuracy: 0.5414

Epoch 11/20

56/56 [=====] - 1s 23ms/step - loss: 0.6651 - accuracy: 0.7584 - val_loss: 1.6330 - val_accuracy: 0.5459

Epoch 12/20

56/56 [=====] - 1s 24ms/step - loss: 0.5668 - accuracy: 0.7891 - val_loss: 1.7141 - val_accuracy: 0.5213

Epoch 13/20

56/56 [=====] - 1s 23ms/step - loss: 0.4977 - accuracy: 0.8265 - val_loss: 1.7319 - val_accuracy: 0.5145

Epoch 14/20

56/56 [=====] - 1s 24ms/step - loss: 0.4099 - accuracy: 0.8493 - val_loss: 1.8707 - val_accuracy: 0.5257

Epoch 15/20

56/56 [=====] - 1s 24ms/step - loss: 0.3986 - accuracy: 0.8677 - val_loss: 1.8504 - val_accuracy: 0.5503

Epoch 16/20

56/56 [=====] - 1s 24ms/step - loss: 0.3467 - accuracy: 0.8705 - val_loss: 1.9885 - val_accuracy: 0.4922

Epoch 17/20

56/56 [=====] - 1s 24ms/step - loss: 0.2689 - accuracy: 0.8973 - val_loss: 2.0118 - val_accuracy: 0.5481


Epoch 18/20

```
56/56 [=====] - 1s 24ms/step - loss: 0.2635 - accur
acy: 0.8934 - val_loss: 2.1038 - val_accuracy: 0.5593
Epoch 19/20
56/56 [=====] - 1s 23ms/step - loss: 0.2285 - accur
acy: 0.9102 - val_loss: 2.1650 - val_accuracy: 0.5436
Epoch 20/20
56/56 [=====] - 1s 24ms/step - loss: 0.1941 - accur
acy: 0.9219 - val_loss: 1.9747 - val_accuracy: 0.5436
```

Visualizing training results

In [119]:

```
1 acc = history.history['accuracy']
2 val_acc = history.history['val_accuracy']
3
4 loss = history.history['loss']
5 val_loss = history.history['val_loss']
6
7 epochs_range = range(epochs)
8
9 plt.figure(figsize=(8, 8))
10 plt.subplot(1, 2, 1)
11 plt.plot(epochs_range, acc, label='Training Accuracy')
12 plt.plot(epochs_range, val_acc, label='Validation Accuracy')
13 plt.legend(loc='lower right')
14 plt.title('Training and Validation Accuracy')
15
16 plt.subplot(1, 2, 2)
17 plt.plot(epochs_range, loss, label='Training Loss')
18 plt.plot(epochs_range, val_loss, label='Validation Loss')
19 plt.legend(loc='upper right')
20 plt.title('Training and Validation Loss')
21 plt.show()
```



Todo: Write your findings after the model fit, see if there is an evidence of model overfit or underfit

Write your findings here


From above graph, training accuracy = 92% but validation accuracy = 54% which is indicating overfitting

In [120]:

```
1 # Todo, after you have analysed the model fit history for presence of underfit or overf
2 data_augmentation = keras.Sequential([
3     layers.experimental.preprocessing.RandomFlip("vertical"),
4     layers.experimental.preprocessing.RandomRotation(10),
5     layers.experimental.preprocessing.RandomZoom(0.1)
6 ])
```

In [121]:

```
1 # Todo, visualize how your augmentation strategy works for one instance of training im
2
3 plt.figure(figsize=(10, 10))
4 for images, labels in train_ds.take(1):
5     for i in range(9):
6         augmented_images = data_augmentation(images)
7         ax = plt.subplot(3, 3, i + 1)
8         plt.imshow(augmented_images[0].numpy().astype("uint8"))
9         plt.axis("off")
```



Todo:

Create the model, compile and train the model

In [122]:

```
1  ## You can use Dropout layer if there is an evidence of overfitting in your findings
2
3  model = Sequential([
4      data_augmentation,
5      layers.experimental.preprocessing.Rescaling(1./255),
6      layers.Conv2D(16, 3, padding='same', activation='relu'),
7      layers.MaxPooling2D(),
8      layers.Conv2D(32, 3, padding='same', activation='relu'),
9      layers.MaxPooling2D(),
10     layers.Conv2D(64, 3, padding='same', activation='relu'),
11     layers.MaxPooling2D(),
12     layers.Dropout(0.2),
13     layers.Flatten(),
14     layers.Dense(128, activation='relu'),
15     layers.Dense(num_classes)
16 ])
```

Compiling the model

In [123]:

```
1  model.compile(optimizer='adam',
2                loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
3                metrics=['accuracy'])
```

Training the model

In [124]:

```
1 ## Your code goes here, note: train your model for 20 epochs
2 epochs = 20
3 history = model.fit(
4     train_ds,
5     validation_data=val_ds,
6     epochs=epochs
7 )
```

Epoch 1/20

56/56 [=====] - 3s 43ms/step - loss: 2.0079 - accuracy: 0.2985 - val_loss: 1.6843 - val_accuracy: 0.4049

Epoch 2/20

56/56 [=====] - 2s 40ms/step - loss: 1.6205 - accuracy: 0.4319 - val_loss: 1.5566 - val_accuracy: 0.4609

Epoch 3/20

56/56 [=====] - 2s 40ms/step - loss: 1.5515 - accuracy: 0.4448 - val_loss: 1.5044 - val_accuracy: 0.4676

Epoch 4/20

56/56 [=====] - 2s 40ms/step - loss: 1.4606 - accuracy: 0.4855 - val_loss: 1.7323 - val_accuracy: 0.4318

Epoch 5/20

56/56 [=====] - 2s 40ms/step - loss: 1.4969 - accuracy: 0.4794 - val_loss: 1.4702 - val_accuracy: 0.4899

Epoch 6/20

56/56 [=====] - 2s 40ms/step - loss: 1.4127 - accuracy: 0.5022 - val_loss: 1.3819 - val_accuracy: 0.5123

Epoch 7/20

56/56 [=====] - 2s 40ms/step - loss: 1.3310 - accuracy: 0.5279 - val_loss: 1.4481 - val_accuracy: 0.5034

Epoch 8/20

56/56 [=====] - 2s 40ms/step - loss: 1.3019 - accuracy: 0.5407 - val_loss: 1.3602 - val_accuracy: 0.5302

Epoch 9/20

56/56 [=====] - 2s 40ms/step - loss: 1.2692 - accuracy: 0.5385 - val_loss: 1.4445 - val_accuracy: 0.5101

Epoch 10/20

56/56 [=====] - 2s 40ms/step - loss: 1.2807 - accuracy: 0.5396 - val_loss: 1.3955 - val_accuracy: 0.5257

Epoch 11/20

56/56 [=====] - 2s 39ms/step - loss: 1.2723 - accuracy: 0.5469 - val_loss: 1.3464 - val_accuracy: 0.5391

Epoch 12/20

56/56 [=====] - 2s 40ms/step - loss: 1.2179 - accuracy: 0.5647 - val_loss: 1.3614 - val_accuracy: 0.5324

Epoch 13/20

56/56 [=====] - 2s 40ms/step - loss: 1.1781 - accuracy: 0.5698 - val_loss: 1.3716 - val_accuracy: 0.5347

Epoch 14/20

56/56 [=====] - 2s 39ms/step - loss: 1.1707 - accuracy: 0.5664 - val_loss: 1.4513 - val_accuracy: 0.4922

Epoch 15/20

56/56 [=====] - 2s 39ms/step - loss: 1.1367 - accuracy: 0.5809 - val_loss: 1.3907 - val_accuracy: 0.5168

Epoch 16/20

56/56 [=====] - 2s 39ms/step - loss: 1.1757 - accuracy: 0.5776 - val_loss: 1.3595 - val_accuracy: 0.5280

Epoch 17/20

56/56 [=====] - 2s 40ms/step - loss: 1.1447 - accuracy: 0.5787 - val_loss: 1.3269 - val_accuracy: 0.5615

Epoch 18/20

56/56 [=====] - 2s 40ms/step - loss: 1.1176 - accuracy: 0.5915 - val_loss: 1.3391 - val_accuracy: 0.5369

Epoch 19/20

56/56 [=====] - 2s 40ms/step - loss: 1.0858 - accuracy: 0.6055 - val_loss: 1.3563 - val_accuracy: 0.5481


Epoch 20/20

56/56 [=====] - 2s 40ms/step - loss: 1.0934 - accuracy: 0.6066 - val_loss: 1.4071 - val_accuracy: 0.5436

Visualizing the results

In [125]:

```
1 acc = history.history['accuracy']
2 val_acc = history.history['val_accuracy']
3
4 loss = history.history['loss']
5 val_loss = history.history['val_loss']
6
7 epochs_range = range(epochs)
8
9 plt.figure(figsize=(8, 8))
10 plt.subplot(1, 2, 1)
11 plt.plot(epochs_range, acc, label='Training Accuracy')
12 plt.plot(epochs_range, val_acc, label='Validation Accuracy')
13 plt.legend(loc='lower right')
14 plt.title('Training and Validation Accuracy')
15
16 plt.subplot(1, 2, 2)
17 plt.plot(epochs_range, loss, label='Training Loss')
18 plt.plot(epochs_range, val_loss, label='Validation Loss')
19 plt.legend(loc='upper right')
20 plt.title('Training and Validation Loss')
21 plt.show()
```



Todo: Write your findings after the model fit, see if there is an evidence of model overfit or underfit. Do you think there is some improvement now as compared to the previous model run?

Comment:

- After using augmentation data there is improvement now as compared to the previous model run.
- Model is still performing poorly on both training and validation data

Todo: Find the distribution of classes in the training dataset.

Context: Many times real life datasets can have class imbalance, one class can have proportionately higher number of samples compared to the others. Class imbalance can have a detrimental effect on the final model quality. Hence as a sanity check it becomes important to check what is the distribution of classes in the data.

In [126]:


```
1  ## Your code goes here.
2  path_list=[]
3  lesion_list=[]
4  for i in class_names:
5
6      for j in data_dir_train.glob(i+'/*.jpg'):
7          path_list.append(str(j))
8          lesion_list.append(i)
9  dataframe_dict_original = dict(zip(path_list, lesion_list))
10 original_df = pd.DataFrame(list(dataframe_dict_original.items()),columns = ['Path','Label'])
11
12 img_df = original_df[['Label']].value_counts().rename_axis('class_name').reset_index(name='count')
13 img_df
14
```

Out[126]:

	class_name	count
0	pigmented benign keratosis	462
1	melanoma	438
2	basal cell carcinoma	376
3	nevus	357
4	squamous cell carcinoma	181
5	vascular lesion	139
6	actinic keratosis	114
7	dermatofibroma	95
8	seborrheic keratosis	77

In [127]:

```
1 import seaborn as sns
2 plt.figure(figsize=(25,10))
3 sns.barplot(x= 'class_name', y = 'count', data=img_df)
4 plt.show()
5
```



Todo: Write your findings here:

- Which class has the least number of samples?
- seborrheic keratosis (77 images)
- Which classes dominate the data in terms proportionate number of samples?
- pigmented benign keratosis(462),melanoma(438),basal cell carcinoma(376),nevus(357)

Todo: Rectify the class imbalance

Context: You can use a python package known as Augmentor (<https://augmentor.readthedocs.io/en/master/>) to add more samples across all classes so that none of the classes have very few samples.

In [128]:

```
1 !pip install Augmentor
```

Looking in indexes: <https://pypi.org/simple>, (<https://pypi.org/simple>,) <https://us-python.pkg.dev/colab-wheels/public/simple/> (<https://us-python.pkg.dev/colab-wheels/public/simple/>)

Requirement already satisfied: Augmentor in /usr/local/lib/python3.7/dist-packages (0.2.10)

Requirement already satisfied: Pillow>=5.2.0 in /usr/local/lib/python3.7/dist-packages (from Augmentor) (7.1.2)

Requirement already satisfied: tqdm>=4.9.0 in /usr/local/lib/python3.7/dist-packages (from Augmentor) (4.64.1)

Requirement already satisfied: numpy>=1.11.0 in /usr/local/lib/python3.7/dist-packages (from Augmentor) (1.21.6)

Requirement already satisfied: future>=0.16.0 in /usr/local/lib/python3.7/dist-packages (from Augmentor) (0.16.0)

To use Augmentor , the following general procedure is followed:

1. Instantiate a Pipeline object pointing to a directory containing your initial image data set.
2. Define a number of operations to perform on this data set using your Pipeline object.
3. Execute these operations by calling the Pipeline's sample() method.

In [129]:

```
1 path_to_training_dataset="Skin cancer ISIC The International Skin Imaging Collaborator
2 import Augmentor
3 for i in class_names:
4     p = Augmentor.Pipeline(path_to_training_dataset + i)
5     p.rotate(probability=0.7, max_left_rotation=10, max_right_rotation=10)
6     p.sample(500) ## We are adding 500 samples per class to make sure that none of the
```

Initialised with 114 image(s) found.

Output directory set to Skin cancer ISIC The International Skin Imaging Collaboration/Train/actinic keratosis/output.

Processing <PIL.Image.Image image mode=RGB size=600x450 at 0x7F90363E0110>: 100%|██████████| 500/500 [00:16<00:00, 29.91 Samples/s]

Initialised with 376 image(s) found.

Output directory set to Skin cancer ISIC The International Skin Imaging Collaboration/Train/basal cell carcinoma/output.

Processing <PIL.JpegImagePlugin.JpegImageFile image mode=RGB size=600x450 at 0x7F9034347F50>: 100%|██████████| 500/500 [00:18<00:00, 27.24 Samples/s]

Initialised with 95 image(s) found.

Output directory set to Skin cancer ISIC The International Skin Imaging Collaboration/Train/dermatofibroma/output.

Processing <PIL.Image.Image image mode=RGB size=600x450 at 0x7F9034304890>: 100%|██████████| 500/500 [00:19<00:00, 25.23 Samples/s]

Initialised with 438 image(s) found.

Output directory set to Skin cancer ISIC The International Skin Imaging Collaboration/Train/melanoma/output.

Processing <PIL.JpegImagePlugin.JpegImageFile image mode=RGB size=962x674 at 0x7F90364C1C90>: 100%|██████████| 500/500 [01:41<00:00, 4.92 Samples/s]

Initialised with 357 image(s) found.

Output directory set to Skin cancer ISIC The International Skin Imaging Collaboration/Train/nevus/output.

Processing <PIL.Image.Image image mode=RGB size=767x576 at 0x7F9034405DD0>: 100%|██████████| 500/500 [01:19<00:00, 6.29 Samples/s]

Initialised with 462 image(s) found.

Output directory set to Skin cancer ISIC The International Skin Imaging Collaboration/Train/pigmented benign keratosis/output.

Processing <PIL.Image.Image image mode=RGB size=600x450 at 0x7F90A050CBD0>: 100%|██████████| 500/500 [00:16<00:00, 30.23 Samples/s]

Initialised with 77 image(s) found.

Output directory set to Skin cancer ISIC The International Skin Imaging Collaboration/Train/seborrheic keratosis/output.

Processing <PIL.JpegImagePlugin.JpegImageFile image mode=RGB size=1024x768 at 0x7F904C096F10>: 100%|██████████| 500/500 [00:38<00:00, 12.97 Samples/s]

Initialised with 181 image(s) found.

Output directory set to Skin cancer ISIC The International Skin Imaging Collaboration/Train/squamous cell carcinoma/output.

Processing <PIL.JpegImagePlugin.JpegImageFile image mode=RGB size=600x450 at 0x7F903433B310>: 100%|██████████| 500/500 [00:16<00:00, 30.99 Samples/s]

Initialised with 139 image(s) found.

Output directory set to Skin cancer ISIC The International Skin Imaging Collaboration/Train/vascular lesion/output.

Processing <PIL.Image.Image image mode=RGB size=600x450 at 0x7F90346FEB50>:
100%|██████████| 500/500 [00:16<00:00, 30.92 Samples/s]

Augmentor has stored the augmented images in the output sub-directory of each of the sub-directories of skin cancer types.. Lets take a look at total count of augmented images.

In [130]:

```
1 image_count_train = len(list(data_dir_train.glob('*/output/*.jpg')))  
2 print(image_count_train)
```

4500

Lets see the distribution of augmented data after adding new images to the original training data.

In [131]:

```
1 path_list_new = [x for x in glob(os.path.join(data_dir_train, '*', 'output', '*.jpg'))]  
2 path_list_new  
3
```

Out[131]:

```
['/content/Skin cancer ISIC The International Skin Imaging Collaboration/T  
rain/dermatofibroma/output/dermatofibroma_original_ISIC_0026417.jpg_6393d1  
19-3431-4b76-8d46-713291129d58.jpg',  
 '/content/Skin cancer ISIC The International Skin Imaging Collaboration/T  
rain/dermatofibroma/output/dermatofibroma_original_ISIC_0030021.jpg_5a40a2  
3f-d1d4-44c6-92d7-74805940b7f2.jpg',  
 '/content/Skin cancer ISIC The International Skin Imaging Collaboration/T  
rain/dermatofibroma/output/dermatofibroma_original_ISIC_0029891.jpg_07c286  
dc-b025-4016-a62d-00cb7e330a9e.jpg',  
 '/content/Skin cancer ISIC The International Skin Imaging Collaboration/T  
rain/dermatofibroma/output/dermatofibroma_original_ISIC_0031271.jpg_685d4c  
36-2404-45b2-aa10-c464b8afb8ed.jpg',  
 '/content/Skin cancer ISIC The International Skin Imaging Collaboration/T  
rain/dermatofibroma/output/dermatofibroma_original_ISIC_0029297.jpg_913fa6  
3a-9e95-49bc-8652-5b3732a76fcf.jpg',  
 '/content/Skin cancer ISIC The International Skin Imaging Collaboration/T  
rain/dermatofibroma/output/dermatofibroma_original_ISIC_0031344.jpg_b27479  
14-f1cb-49be-b4ba-c14b034e8760.jpg']
```

In [132]:

```
1 lesion_list_new = [os.path.basename(os.path.dirname(os.path.dirname(y))) for y in glob(
2 lesion_list_new
```

Out[132]:

[illegible]

In [133]:

```
1 dataframe dict new = dict(zip(path list new, lesion list new))
```

In [134]:

```
1 df2 = pd.DataFrame(list(dataframe_dict_new.items()), columns = ['Path', 'Label'])
2 new df = original df.append(df2)
```

In [135]:

```
1 new_df['Label'].value_counts()
```

Out[135]:

```
pigmented benign keratosis      962
melanoma                        938
basal cell carcinoma            876
nevus                           857
squamous cell carcinoma         681
vascular lesion                 639
actinic keratosis               614
dermatofibroma                 595
seborrheic keratosis           577
Name: Label, dtype: int64
```

So, now we have added 500 images to all the classes to maintain some class balance. We can add more images as we want to improve training process.

Todo: Train the model on the data created using Augmentor

In [136]:

```
1 batch_size = 32
2 img_height = 180
3 img_width = 180
```

Todo: Create a training dataset

In [137]:

```
1 train_ds = tf.keras.preprocessing.image_dataset_from_directory(
2     data_dir_train,
3     seed=123,
4     validation_split = 0.2,
5     subset = "training",
6     image_size=(img_height, img_width),
7     batch_size=batch_size)
```

Found 6739 files belonging to 9 classes.
Using 5392 files for training.

Todo: Create a validation dataset

In [138]:

```
1 val_ds = tf.keras.preprocessing.image_dataset_from_directory(
2     data_dir_train,
3     seed=123,
4     validation_split = 0.2,
5     subset = 'validation', ## Todo choose the correct parameter value, so that only valid
6     image_size=(img_height, img_width),
7     batch_size=batch_size)
```

Found 6739 files belonging to 9 classes.
Using 1347 files for validation.

Todo: Create your model (make sure to include normalization)

In [139]:

```
1 ## your code goes here
2 AUTOTUNE = tf.data.experimental.AUTOTUNE
3
4 train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
5 val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
```

Todo: Compile your model (Choose optimizer and loss function appropriately)

In [140]:

```
1  ## your code goes here
2  model = Sequential([
3      layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
4
5      layers.BatchNormalization(),
6      layers.Conv2D(128, 3, padding = 'same', activation='relu'),
7      layers.MaxPooling2D(),
8
9      layers.BatchNormalization(),
10     layers.Conv2D(256, 3, padding = 'same', activation='relu'),
11     layers.MaxPooling2D(),
12
13     layers.BatchNormalization(),
14     layers.Conv2D(512, 3, padding = 'same', activation='relu'),
15     layers.MaxPooling2D(),
16     layers.Dropout(0.4),
17
18     layers.Flatten(),
19
20     layers.BatchNormalization(),
21     layers.Dense(128, activation='relu'),
22
23     layers.BatchNormalization(),
24     layers.Dense(32, activation='relu'),
25     layers.Dropout(0.4),
26
27     layers.Dense(num_classes)
28 ])
```

In [141]:

```
1  model.compile(optimizer='adam',
2                loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
3                metrics=['accuracy'])
```

Todo: Train your model

In [142]:

```
1 epochs = 30
2 history = model.fit(
3     train_ds,
4     validation_data=val_ds,
5     epochs=epochs
6 )
```

Epoch 1/30

169/169 [=====] - 64s 229ms/step - loss: 1.8820 - accuracy: 0.3353 - val_loss: 3.2016 - val_accuracy: 0.1158

Epoch 2/30

169/169 [=====] - 36s 211ms/step - loss: 1.4632 - accuracy: 0.4638 - val_loss: 1.9519 - val_accuracy: 0.2814

Epoch 3/30

169/169 [=====] - 36s 215ms/step - loss: 1.2724 - accuracy: 0.5336 - val_loss: 1.1598 - val_accuracy: 0.5754

Epoch 4/30

169/169 [=====] - 36s 213ms/step - loss: 1.1033 - accuracy: 0.5938 - val_loss: 0.9303 - val_accuracy: 0.6763

Epoch 5/30

169/169 [=====] - 36s 213ms/step - loss: 0.9662 - accuracy: 0.6465 - val_loss: 0.8649 - val_accuracy: 0.6793

Epoch 6/30

169/169 [=====] - 36s 214ms/step - loss: 0.8397 - accuracy: 0.6981 - val_loss: 0.7615 - val_accuracy: 0.7253

Epoch 7/30

169/169 [=====] - 36s 214ms/step - loss: 0.7323 - accuracy: 0.7329 - val_loss: 0.7668 - val_accuracy: 0.7127

Epoch 8/30

169/169 [=====] - 36s 214ms/step - loss: 0.6291 - accuracy: 0.7674 - val_loss: 0.7691 - val_accuracy: 0.7298

Epoch 9/30

169/169 [=====] - 36s 214ms/step - loss: 0.5639 - accuracy: 0.7921 - val_loss: 0.7363 - val_accuracy: 0.7439

Epoch 10/30

169/169 [=====] - 36s 214ms/step - loss: 0.4873 - accuracy: 0.8201 - val_loss: 0.5703 - val_accuracy: 0.8018

Epoch 11/30

169/169 [=====] - 36s 214ms/step - loss: 0.4377 - accuracy: 0.8344 - val_loss: 0.5849 - val_accuracy: 0.8077

Epoch 12/30

169/169 [=====] - 36s 214ms/step - loss: 0.3605 - accuracy: 0.8655 - val_loss: 0.5746 - val_accuracy: 0.8070

Epoch 13/30

169/169 [=====] - 36s 214ms/step - loss: 0.3495 - accuracy: 0.8731 - val_loss: 0.7671 - val_accuracy: 0.7372

Epoch 14/30

169/169 [=====] - 36s 213ms/step - loss: 0.3547 - accuracy: 0.8624 - val_loss: 0.6478 - val_accuracy: 0.8092

Epoch 15/30

169/169 [=====] - 36s 214ms/step - loss: 0.2910 - accuracy: 0.8906 - val_loss: 0.6231 - val_accuracy: 0.8085

Epoch 16/30

169/169 [=====] - 36s 214ms/step - loss: 0.2481 - accuracy: 0.9102 - val_loss: 0.5332 - val_accuracy: 0.8263

Epoch 17/30

169/169 [=====] - 36s 213ms/step - loss: 0.2446 - accuracy: 0.9101 - val_loss: 0.6931 - val_accuracy: 0.8129


Epoch 18/30

169/169 [=====] - 36s 213ms/step - loss: 0.2648 - accuracy: 0.8986 - val_loss: 0.6191 - val_accuracy: 0.8151
Epoch 19/30
169/169 [=====] - 36s 214ms/step - loss: 0.2519 - accuracy: 0.9049 - val_loss: 0.6489 - val_accuracy: 0.8070
Epoch 20/30
169/169 [=====] - 36s 214ms/step - loss: 0.2282 - accuracy: 0.9177 - val_loss: 0.7648 - val_accuracy: 0.8099
Epoch 21/30
169/169 [=====] - 36s 213ms/step - loss: 0.2057 - accuracy: 0.9269 - val_loss: 0.7490 - val_accuracy: 0.8070
Epoch 22/30
169/169 [=====] - 36s 213ms/step - loss: 0.2488 - accuracy: 0.9114 - val_loss: 0.7241 - val_accuracy: 0.8010
Epoch 23/30
169/169 [=====] - 36s 213ms/step - loss: 0.2001 - accuracy: 0.9251 - val_loss: 0.7374 - val_accuracy: 0.8203
Epoch 24/30
169/169 [=====] - 36s 214ms/step - loss: 0.1965 - accuracy: 0.9288 - val_loss: 0.7784 - val_accuracy: 0.8151
Epoch 25/30
169/169 [=====] - 36s 213ms/step - loss: 0.1706 - accuracy: 0.9358 - val_loss: 0.6529 - val_accuracy: 0.8322
Epoch 26/30
169/169 [=====] - 36s 213ms/step - loss: 0.1599 - accuracy: 0.9420 - val_loss: 0.7501 - val_accuracy: 0.8137
Epoch 27/30
169/169 [=====] - 36s 214ms/step - loss: 0.1391 - accuracy: 0.9514 - val_loss: 0.6406 - val_accuracy: 0.8359
Epoch 28/30
169/169 [=====] - 36s 213ms/step - loss: 0.1233 - accuracy: 0.9562 - val_loss: 0.7358 - val_accuracy: 0.8337
Epoch 29/30
169/169 [=====] - 36s 213ms/step - loss: 0.1315 - accuracy: 0.9516 - val_loss: 0.8800 - val_accuracy: 0.8055
Epoch 30/30
169/169 [=====] - 36s 214ms/step - loss: 0.1355 - accuracy: 0.9514 - val_loss: 0.8075 - val_accuracy: 0.8048

Todo: Visualize the model results

In [143]:

```
1 acc = history.history['accuracy']
2 val_acc = history.history['val_accuracy']
3
4 loss = history.history['loss']
5 val_loss = history.history['val_loss']
6
7 epochs_range = range(epochs)
8
9 plt.figure(figsize=(8, 8))
10 plt.subplot(1, 2, 1)
11 plt.plot(epochs_range, acc, label='Training Accuracy')
12 plt.plot(epochs_range, val_acc, label='Validation Accuracy')
13 plt.legend(loc='lower right')
14 plt.title('Training and Validation Accuracy')
15
16 plt.subplot(1, 2, 2)
17 plt.plot(epochs_range, loss, label='Training Loss')
18 plt.plot(epochs_range, val_loss, label='Validation Loss')
19 plt.legend(loc='upper right')
20 plt.title('Training and Validation Loss')
21 plt.show()
```



Todo: Analyze your results here. Did you get rid of underfitting/overfitting? Did class rebalance help?

Comments

In the final model, there is no sign of underfitting/overfitting.

Class rebalanced improved the model performance on both training and validation data.

In [144]:

```
1 print("----- Accuracy -----")
2
3 print("Accuracy          : ", acc[-1])
4 print("Validation Accuracy : ", val_acc[-1])
5 print("Loss              : ", loss[-1])
6 print("Validation Loss     : ", val_loss[-1])
```

```
----- Accuracy -----
Accuracy          : 0.9514095187187195
Validation Accuracy : 0.8047512769699097
Loss              : 0.13551045954227448
Validation Loss     : 0.8074936866760254
```

In [145]:

```
1 from keras.utils.vis_utils import plot_model
2 plot_model(model, to_file='model_plot.png', show_shapes=True, show_layer_names=True)
```

Out[145]:

